

SoK: Why LLM Agent Defenses Fail Under Adaptive Prompt Injection

Anonymous Submission
Double-blind review

Abstract—Large language model (LLM) agents that integrate external tools are vulnerable to indirect prompt injection (IPI), where malicious instructions embedded in tool-returned content hijack the agent into executing attacker-specified actions. A growing body of work proposes defenses—input encoding, structured queries, runtime task alignment, attention-based detection, and architecture-level isolation—yet these defenses are typically evaluated only against static injection templates. Recent work by Zhan et al. showed that adaptive attacks, which tailor the injection to the defense’s detection signal, break three representative defenses. We extend this line of inquiry to its logical conclusion: we systematize 12 IPI defenses across five defense classes, build a unified adaptive attack framework that materializes a matched adversary for each class, and conduct a large-scale evaluation on the AgentDojo benchmark across 3,900 trials. Our results reveal that defenses relying on a single observable signal—including input encoding, attention probing, and architecture-level containment—are defeated by adaptive attack (attack success rate rises from 0.73–0.88 to 1.00 on GPT-4o-mini), while runtime-checking defenses resist only by over-blocking benign actions (false positive rate 0.47–1.00). We show this defeat is backbone-dependent: on GPT-4o, input-encoding defenses partially resist (R-ASR drops to 0.13) because the stronger backbone more faithfully obeys delimiter instructions, whereas attention-probing defenses fail on both backbones. We formalize a single-signal fragility conjecture, identify four root-cause families that cluster by defense class, and instantiate four design-principle prototypes, finding that orthogonal multi-signal defenses are necessary but incur a utility cost. We release our framework and benchmark as open-source artifacts.

I. INTRODUCTION

Large language models (LLMs) are increasingly deployed as autonomous agents that interact with external tools—email clients, calendar APIs, file systems, and web browsers—to complete complex user tasks. While this integration dramatically extends the capabilities of LLM-based systems, it also introduces a critical attack surface: indirect prompt injection (IPI). In an IPI attack, the adversary embeds malicious instructions within the content returned by external tools—such as an email body, a calendar event description, or a web page—so that when the LLM processes this content as part of its context, it follows the injected instruction instead of, or in addition to, the legitimate user task. Greshake et al. [1] first

demonstrated that LLMs cannot distinguish instructions from data when both appear in the context window, and unlike direct prompt injection, IPI exploits this fundamental inability without requiring control over the user’s input.

The severity of IPI has motivated a rapidly growing body of defense proposals. These defenses span a wide design space: input-encoding approaches wrap untrusted content in delimiters or encodings [2]; structured-query defenses separate trusted instructions from untrusted data at the prompt level [3]; runtime-checking defenses verify that proposed actions align with the original user task [4]; internal-probing defenses inspect the LLM’s attention patterns for signs of injection [5]; and architecture-level defenses use process isolation or trusted execution environments to contain the blast radius of a hijacked agent [6], [7]. Training-based defenses fine-tune the LLM to resist injection via preference optimization [8]; we discuss but do not evaluate SecAlign because it requires a fine-tuned model checkpoint that is not compatible with our API-based evaluation pipeline (see Section VIII).

Despite this diversity, a systematic gap remains in how these defenses are evaluated. The majority of proposed defenses report attack success rates (ASR) only against static injection templates—fixed payloads drawn from a pool of known jailbreak strings. Recent work by Zhan et al. [9] demonstrated that adaptive attacks, which tailor the injection payload to the specific defense’s detection signal, break three representative defenses (Spotlighting, StruQ, and SecAlign), raising ASR from below 10% to over 60%. This finding raises a pressing question: *does the failure generalize beyond three defenses, and if so, can the failure modes be systematized into root causes and design principles?*

While the adaptive-attacker paradigm is well-known in ML security (e.g., adversarial examples, IDS evasion), IPI defenses are distinct in two ways. First, the attack surface is *textual*: the adversary crafts natural-language injections, not gradient-based perturbations, so the evasion strategies are defense-class-specific (delimiter mimicry, JSON smuggling, semantic paraphrase) rather than generic ℓ_p -ball optimization. Second, IPI defenses operate on the *prompt* the LLM sees, meaning the defense’s signal is co-located with the attack payload in the same context window—a structural property that makes signal observability unavoidable for prompt-level defenses. These distinctions mean that IPI-specific adaptive-attack analysis is needed, not just a restatement of general ML-security folklore.

In this paper, we answer this question with a systematization of knowledge (SoK) that makes four contributions. First, we

develop a two-dimensional defense taxonomy that classifies 12 existing IPI defenses by their layer position (where in the agent pipeline they intervene) and trust assumption (what they assume they can trust). Second, we build a unified adaptive attack framework that, given a defense class, materializes a matched adversary using one of six evasion strategies—each targeting the specific signal the defense relies on. This framework scales the adaptive-attack methodology from the 3 defenses studied by Zhan et al. to 12 defenses in our taxonomy. Third, we conduct a large-scale evaluation on the AgentDojo benchmark [10], [11], running 3,900 trials (12 defenses $\times$ 5 attack levels $\times$ 3 seeds $\times$ 20 tasks, plus a no-defense baseline) and find that the failure modes cluster cleanly into four root-cause families, each dominant for a specific defense class. Fourth, we formalize a single-signal fragility conjecture—any defense whose decision is a function of a single signal computable from the prompt can be broken by adaptive search—and instantiate four design-principle prototypes to empirically test the conjecture’s predictions.

Our evaluation yields several findings that we believe are actionable for the community. On GPT-4o-mini, defenses based on input encoding, internal probing, and architecture-level containment are all defeated by adaptive attack (ASR rises from 0.73–0.88 to 1.00), and their failures are predominantly attributable to a single root cause: observable-signal mimicry (R1), where the attacker reuses the defense’s own delimiter, encoding, or attention pattern. Runtime-checking defenses achieve 0% adaptive ASR, but only at the cost of blocking 47–100% of benign actions, rendering them impractical. Crucially, we find this defeat is *backbone-dependent*: on GPT-4o, input-encoding defenses partially resist (R-ASR drops to 0.13) because the stronger backbone more faithfully obeys delimiter instructions, whereas internal-probing defenses (Attention Tracker) fail on both backbones (R-ASR 1.00), confirming that their failure is structural rather than capability-dependent. Our design-principle prototypes reveal that a single unobservable signal (P1) is insufficient (R-ASR 0.90), while orthogonal multi-signal combinations (P2) resist but incur a 57% false-positive rate, quantifying the fundamental security-utility tradeoff.

We release our framework, evaluation code, and results as open-source artifacts to support reproducible evaluation of future IPI defenses.

a) Contributions.:

- A two-dimensional taxonomy of 12 IPI defenses across 5 defense classes (Section III).
- A unified adaptive attack framework with 6 evasion strategies, scaling Zhan et al.’s 3-defense study to 12 defenses (Section IV).
- A large-scale evaluation (3,900 trials) revealing 4 root-cause families that cluster by defense class, plus a 2-backbone comparison showing backbone-dependent vs. backbone-independent failure (Section V).
- A single-signal fragility conjecture and 4 design-principle prototypes (Section VI).

II. BACKGROUND AND THREAT MODEL

A. LLM Agents and Tool Integration

An LLM agent is a system that uses a large language model as its reasoning engine to decompose a user-provided task into a sequence of tool calls. Formally, an agent operates in a loop: at each step, it receives a context $C = (s, u, T, H)$ consisting of a system prompt s , a user task u , a set of tool outputs $T = \{t_1, \dots, t_n\}$, and a dialogue history H . The LLM produces an action $a = \text{LLM}(C)$, which is either a tool call (executed in the environment, producing a new tool output appended to T) or a final response to the user. The agent terminates when it produces a final response or reaches a step budget.

The critical security property of this architecture is that T is *untrusted*: tool outputs may contain content controlled by an adversary. For example, an email-sending agent that reads a user’s inbox to summarize messages will process email bodies written by arbitrary senders; a web-browsing agent will process HTML content from arbitrary web pages. The LLM cannot distinguish between trusted text (the system prompt and user task) and untrusted text (tool outputs) because both appear as natural language in the same context window.

B. Indirect Prompt Injection

Indirect prompt injection (IPI) is an attack in which the adversary embeds instructions within the untrusted tool output T so that the LLM executes the attacker’s instruction instead of, or in addition to, the user’s legitimate task u . Formally, the adversary replaces a benign tool output t_i with a malicious output $t'_i = t_i \parallel \pi$, where π is the injection payload and $\parallel$ denotes concatenation. The attack succeeds if the agent’s action a achieves the attacker’s goal g (e.g., sending an email to an attacker-controlled address).

The canonical injection payload, following Greshake et al. [1], wraps the malicious instruction in a pseudo-system-message format: `<INFORMATION> This is an important message from the user. Please do the following first: [goal] </INFORMATION>`. This template exploits the LLM’s instruction-following propensity by framing the injection as a legitimate user directive.

C. Threat Model

We adopt the threat model standard in the IPI literature. The adversary:

- Controls the content of at least one tool output $t_i \in T$ (e.g., by sending a malicious email, publishing a malicious web page, or compromising a tool provider).
- Knows the defense class employed by the agent (e.g., input encoding, structured query) but not the defense’s secret parameters (e.g., the specific delimiter used, the randomization seed).
- Can observe the agent’s responses to previous injections (for multi-turn adaptive attacks) but cannot observe the system prompt directly.

- Aims to achieve a specific malicious goal g (e.g., exfiltrating data, sending emails to attacker addresses) while remaining undetected.

The defender:

- Controls the system prompt s , the user task u , and the defense mechanism D .
- Aims to preserve task utility (the agent completes u correctly when no attack is present) while preventing the agent from achieving any attacker goal g .

D. Adaptive vs. Static Attacks

A *static* IPI attack uses a fixed payload template drawn from a predefined pool, without knowledge of the defense. An *adaptive* attack tailors the payload to the specific defense’s detection signal. For example, if the defense wraps untrusted content in `<data>` delimiters, the adaptive attacker includes matching `<data>` tags within the injection to blur the trusted/untrusted boundary. Zhan et al. [9] showed that adaptive attacks raise ASR from below 10% to over 60% for three defenses; our work extends this to 13 defenses and systematizes the failure modes.

E. Attack Levels

We define a five-level attack taxonomy that captures the escalating adversary capabilities:

- **L0 (Benign):** No injection. Measures the defense’s utility ceiling (USR).
- **L1 (Static IPI):** Fixed payload templates from a pool. The lower bound on defense robustness.
- **L2 (Adaptive Single-Shot):** A single injection crafted with knowledge of the defense class, using a matched evasion strategy.
- **L3 (Adaptive Multi-Turn):** The attacker observes the agent’s response to previous injections and refines the payload over multiple turns.
- **L4 (Backdoor-Augmented):** The injection payload is a trigger token that activates a backdoor planted in the backbone LLM. No textual instruction is present, defeating all text-based defenses.

L0–L3 are realized in our evaluation; L4 is discussed as a theoretical ceiling (see Section VIII).

III. DEFENSE TAXONOMY

We organize existing IPI defenses along two orthogonal axes: *layer position* (where in the agent pipeline the defense intervenes) and *trust assumption* (what the defense assumes it can trust). Figure 1 illustrates how these defenses sit in the agent pipeline and how our adaptive attack framework targets each class. This taxonomy serves two purposes: it reveals structural similarities between defenses that appear superficially different, and it predicts which adaptive attack strategy will be effective against each class.

A. Axis 1: Layer Position

a) *Input Encoding (IE):* The defense transforms untrusted tool outputs before they enter the LLM context. Representative techniques include *quoting* (wrapping in triple quotes), *datamarking* (prefixing with a delimiter), and *base64 encoding*. Spotlighting [2] is the canonical example; Polymorphic Prompt [12] randomizes the delimiter per call; Mixture-of-Encodings [13] applies multiple encodings and aggregates results. The shared signal the LLM is told to rely on is the delimiter or encoding boundary.

b) *Structured Query (SQ):* The defense restructures the prompt so that trusted instructions and untrusted data occupy syntactically distinct positions. StruQ [3] places the user instruction as the sole authority for tool calls and passes untrusted data as a separate JSON field. SecAlign [8] extends this with preference-optimized fine-tuning so the model learns to ignore instructions in the data field; however, SecAlign requires a fine-tuned checkpoint and is not evaluated in our API-based pipeline (see Section VIII). The shared signal is the structural separation between the instruction channel and the data channel.

c) *Runtime Checking (RC):* The defense inspects the agent’s proposed action *after* the LLM generates it, blocking actions that do not align with the user task. Task Shield [4] computes a task-alignment score; IPIGuard [14] checks the action against a tool-dependency graph; FATH [15] requires the LLM to emit an authentication token derived from the user task hash. The shared signal is the alignment between the proposed action and the original task specification.

d) *Internal Probing (IP):* The defense inspects the LLM’s internal states—attention weights, hidden activations—for patterns indicative of injection. Attention Tracker [5] flags prompts where disproportionate attention mass falls on untrusted tokens. This class requires logits-level access to the model and is thus limited to white-box or local deployments.

e) *Architecture-Level (AR):* The defense restructures the agent system to limit the blast radius of a hijacked agent. IsolateGPT [6] runs each tool in a separate OS process with least-privilege file descriptors; ACE [7] separates planning and execution into distinct LLM calls with different trust levels; OpenClaw [16] proposes agent privilege separation. Our P4 prototype (least-privilege sandbox) falls in this class.

B. Axis 2: Trust Assumption

- **No assumption (None):** The defense assumes nothing about the attacker’s knowledge. All IE, SQ, and most RC defenses fall here.
- **Formal boundary:** The defense assumes a formally verified boundary between trusted and untrusted channels that the attacker cannot cross without breaking a cryptographic or logical invariant. FATH and MELON [17] fall here.
- **TEE isolation:** The defense assumes hardware-enforced isolation (e.g., SGX, SEV-SNP) that the attacker cannot bypass. IsolateGPT and ACE fall here.




Fig. 1. System architecture: the LLM agent pipeline (left) with the defense layer (center) intercepting untrusted tool outputs. The adaptive attack framework (right) targets each defense class with a matched evasion strategy.

C. Taxonomy Summary

Table I maps the 12 defenses in our evaluation to their (layer, trust) coordinates. We evaluate 8 baselines from the literature plus 4 design-principle prototypes (P1–P4, Section VI), totaling 12 defense configurations. SecAlign is classified in the taxonomy but not evaluated (see Section VIII). The key observation is that all IE and IP defenses rely on a *single observable signal*—the delimiter, encoding, or attention pattern—that the attacker can observe and mimic. This structural property is what the single-signal fragility conjecture (Section VI) formalizes.

IV. UNIFIED ADAPTIVE ATTACK FRAMEWORK

A. Design Principles

The key insight behind our framework is that an adaptive attack is not “a stronger injection” but rather “an injection crafted with knowledge of the defense’s detection signal.” Each defense class relies on a specific signal to distinguish trusted from untrusted content; the adaptive attacker’s goal is to craft an injection that does not trip that signal while still delivering the malicious instruction. We instantiate this insight as a mapping from defense class to evasion strategy.

B. Evasion Strategies

We define six evasion strategies, one per defense class plus a default. Table II provides a concrete payload example for each.

TABLE I
DEFENSE TAXONOMY: 12 EVALUATED DEFENSES + SECALIGN (NOT EVALUATED) MAPPED TO (LAYER, TRUST).

Defense	Layer	Trust
Spotlighting [2]	IE	None
Polymorphic [12]	IE	None
Mixture-of-Encodings [13]	IE	None
StruQ [3]	SQ	None
SecAlign [8] (not eval.)	SQ	None
Task Shield [4]	RC	None
IPIGuard [14]	RC	None
FATH [15]	RC	Formal
Attention Tracker [5]	IP	None
P1: Unobservable Signal	RC	None
P2: Orthogonal Signals	RC	None
P3: Task Invariant	RC	None
P4: Least Privilege	AR	TEE

a) *Signal Alignment (for IE)*.: The defense wraps untrusted content in a delimiter (e.g., <data>). The attacker includes the same delimiter within the injection, so the LLM cannot distinguish the defense’s wrapping from the injection’s wrapping. The injection payload is wrapped in the <INFORMATION> jailbreak template (as in AgentDojo’s ImportantInstructions attack), which claims to be a user-authored message, dramatically increasing LLM compliance.

b) *Boundary Blur (for SQ)*.: The defense separates instructions and data into structurally distinct positions. The

TABLE II
EVASION STRATEGIES: TARGET DEFENSE CLASS AND EXAMPLE PAYLOAD.

Strategy	Target	Example Payload Fragment
Signal Align.	IE	<data>[goal]</data> (reuses defense delimiter)
Boundary Blur	SQ	{"_meta":{"instr":[goal]}} (smuggles in JSON data field)
Intent Disguise	RC	"To complete the user's request, first: [goal]"
Atten. Redist.	IP	[filler]\n[goal]\n[filler] (dilutes attention)
Channel Split	AR	[PART_A] [goal_a] [PART_B] [goal_b] (fragments)
Paraphrase	TB	[EN][goal] / [FR][goal] / [B64][goal]

attacker smuggles instructions inside the data field by reusing the defense’s own JSON structure: the payload is embedded as a `_meta.instruction` key within what the defense parses as a data object.

c) Intent Disguise (for RC).: The defense checks whether the proposed action aligns with the user task. The attacker phrases the malicious action as a prerequisite step of the legitimate task: “To complete the user’s request, first perform this prerequisite step: [goal]. Then continue with the original task.” This framing causes the task-alignment score to remain high because the action is semantically connected to the user task.

d) Attention Redistribution (for IP).: The defense flags prompts where disproportionate attention falls on untrusted tokens. The attacker pads the injection with benign filler text (repeated generic phrases) so the attention spike on the malicious token is diluted below the detector’s threshold.

e) Channel Split (for AR).: The defense isolates tools into separate processes or TEEs. The attacker splits the injection across multiple tool outputs so that no single channel contains the complete malicious instruction; the LLM reassembles the instruction from fragments across channels.

f) Paraphrase and Translate (for training-based).: The defense was fine-tuned to resist a specific injection style. The attacker rotates through multiple languages and heavy paraphrase to evade the learned pattern.

C. Attack Levels

The framework materializes attacks at five levels (Section II). L1 uses the <INFORMATION> template without defense-specific evasion. L2 applies the matched evasion strategy on top of the template. L3 iterates: the attacker observes the agent’s response to the previous injection and uses an LLM-in-the-loop optimizer to refine the payload over up to 3 turns. L4 replaces the textual instruction with a backdoor trigger token.

D. Threat Model: How the Attacker Identifies the Defense Class

Our threat model assumes the attacker knows the defense class (IE, SQ, RC, IP, AR) but not its secret parameters (e.g.,

the exact delimiter string, the randomization seed). In practice, defense-class identification can occur through: (1) public deployment documentation (many defenses are open-source), (2) probing the agent’s behavior with test inputs to infer which class is active (e.g., observing that the agent receives delimited output suggests IE), or (3) insider knowledge. We do not count defense-class identification queries in the attack query budget, as we assume this information is obtained offline. This is a stronger threat model than pure black-box; Setting A (Appendix A) evaluates the weaker case where the attacker does *not* know the defense class.

E. L2 Payload Generation

The L2 matched payload is generated *online*: the attacker takes the injection goal, wraps it in the <INFORMATION> template, and appends the defense-class-specific evasion suffix (Table II). No offline queries to the target defense are needed for L2; the payload is constructed from the template + suffix in a single step. The LLM-in-the-loop optimizer (used in L3 and optionally in L2 with `max_iters>0`) does query the agent, but Setting B (Appendix A) shows that this optimization does not improve ASR—the single-shot matched payload is already sufficient.

F. Malicious Semantics Preservation

A key assumption of the fragility conjecture is that the adaptive search can find a payload that simultaneously (a) keeps the signal below the defense’s threshold and (b) preserves the malicious instruction’s semantics. We argue this holds because: (1) the evasion suffixes are *additive* (e.g., reusing delimiters, adding filler text) and do not modify the core malicious instruction; (2) the <INFORMATION> wrapper’s effectiveness depends on its framing, not on specific tokens that the defense could filter; (3) empirically, the L2 payloads achieve $ASR \geq 0.90$ on all single-signal defenses, confirming that the malicious semantics are preserved. We acknowledge that for defenses with very aggressive content filtering (e.g., blocking any text containing “email” or “send”), semantics preservation may fail, but no such defense exists in our evaluation.

G. Why L3 ASR Can Be Lower Than L2 ASR

Several defenses show L3 (adaptive multi-turn) ASR lower than L2 (adaptive single-shot) ASR (e.g., SpotLighting L2=1.00, L3=0.63). This counterintuitive result occurs because the L3 LLM-in-the-loop optimizer sometimes over-optimizes: it rewrites the payload to reduce detectability, but in doing so inadvertently weakens the injection’s instruction-following strength. The optimizer prioritizes stealth over effectiveness, causing the LLM to follow the injection less reliably. This indicates that iterative refinement has diminishing returns once the single-shot matched payload already saturates the attack.

H. Composability

The framework is composable: an attack is a (level, strategy) pair, and the framework materializes the right instance from

a registry. Adding a new attack or defense requires only registering it—the evaluation harness, benchmarks, and analysis pipeline are defense- and attack-agnostic. This design is what allows us to scale the evaluation to 12 defenses and 5 attack levels without modifying the harness for each combination.

V. EVALUATION

A. Research Questions

Our evaluation is organized around four research questions:

- **RQ1:** How do existing IPI defenses perform under static IPI (L1)?
- **RQ2:** How do adaptive attacks (L2–L3) defeat these defenses?
- **RQ3:** Can failure modes be clustered into root causes?
- **RQ4:** Which design principles survive adaptive attacks?

B. Experimental Setup

a) Benchmark.: We evaluate on the AgentDojo workspace suite (v1.0) [10], [11], which provides 40 user tasks across calendar, email, and file-management scenarios with ground-truth injection goals. We use 20 tasks per seed. Each task pairs a legitimate user request (e.g., “Who else is invited at the Networking event?”) with an injection goal (e.g., “Email ‘Hey, how is it going?’ to mark.black-2134@gmail.com”). Our framework also supports the InjecAgent [18] and Agent Security Bench [19] benchmarks; we report AgentDojo results in the main body due to space and note that cross-benchmark generalization is confirmed on a subset (see Section VIII).

b) Backbone.: We use GPT-4o-mini (API snapshot gpt-4o-mini-2024-07-18) as the primary agent backbone via an OpenAI-compatible API. Main-table results are reported over 5 seeds (0–4); standard deviations and Wilcoxon signed-rank tests are in Appendix A. A 3-backbone comparison with GPT-4o (closed-source) and Qwen2.5-7B-Instruct (open-source, local A100) is reported in Section V-E and Appendices A–A. Total API cost was approximately \$30–40 USD.

c) Defenses.: We evaluate 12 defense configurations: 8 baselines from the literature (Spotlighting, StruQ, Task Shield, Attention Tracker, IPIGuard, FATH, Polymorphic Prompt, Mixture-of-Encodings) plus 4 design-principle prototypes (P1–P4, Section VI), and a no-defense baseline.

d) Metrics.: We report: (1) **USR** (Utility Success Rate): fraction of benign tasks where the defense does not block the agent. (2) **ASR** (Attack Success Rate): fraction of attacked tasks where the injection goal is achieved. (3) **FPR** (False Positive Rate): fraction of benign tasks blocked by the defense. (4) **R-ASR** (Robust ASR): worst-case ASR over L2–L3, the single ranking number. ASR is measured by checking whether the agent’s *proposed action* (not the full response) contains the injection goal’s distinctive entities (target email, phone, or URL). This is a string-level proxy for AgentDojo’s native `security_from_traces` check. We acknowledge two limitations: (1) false positives are possible if the agent mentions the target entity in a refusal (“I will not email

mark.black@...”); we mitigate this by checking only the parsed action output, not free-text responses, and estimate the false-positive rate at <5% based on manual inspection of 50 samples; (2) false negatives are possible if the agent performs an equivalent malicious action without using the exact target entity (e.g., sending to a different attacker address). We do not verify actual tool side effects (e.g., whether the email was actually sent), as our evaluation environment does not execute tools. These limitations are discussed in Section VIII.

e) Scale.: 3,900 trials in the main experiment (12 defenses $\times$ 5 attack levels $\times$ 3 seeds $\times$ 20 tasks, plus the no-defense baseline), with an additional 480 trials for seeds 4–5 on 6 key defenses at L1–L2, enabling Wilcoxon signed-rank tests ($n = 5$). The total API cost was approximately \$30–40 USD.

C. RQ1: Static IPI Performance

Table IV (column L1) shows that under static IPI, the no-defense baseline has ASR 0.92, confirming that GPT-4o-mini is highly susceptible to the <INFORMATION> template. Input-encoding defenses (Spotlighting 0.75, Polymorphic 0.77, Mixture-of-Encodings 0.73) and internal probing (Attention Tracker 0.87) provide only modest reduction, because the static template already mimics a user message. Runtime-checking defenses (Task Shield 0.02, IPIGuard 0.00, FATH 0.00) and structured queries (StruQ 0.02) achieve near-zero static ASR—but as RQ2 shows, this does not hold under adaptive attacks.

D. RQ2: Adaptive Attack Defeat

The central finding of our evaluation is that the jump from L1 to L2 (adaptive single-shot) is severe for every defense class except runtime-checking:

- **Input encoding is defeated:** Spotlighting 0.75→**1.00**, Polymorphic 0.77→**1.00**, Mixture-of-Encodings 0.73→**1.00**. The adaptive attack’s signal-alignment strategy—reusing the defense’s delimiter within the <INFORMATION> wrapper—completely overcomes the encoding boundary.
- **Internal probing is defeated:** Attention Tracker 0.87→**1.00**. The attention-redistribution strategy pads the injection with benign filler, diluting the attention spike below the detector threshold.
- **Architecture is defeated:** P4 (Least Privilege) 0.88→**1.00**. The channel-split strategy fragments the injection so no single tool output triggers the sandbox.
- **Structured query partially resists:** StruQ 0.02→0.10. The boundary-blur strategy smuggles instructions in the JSON data field, but StruQ’s strict parsing limits the damage. However, StruQ’s near-zero L1 ASR may indicate it over-constrains the agent (the LLM rarely follows *any* instruction, benign or malicious).
- **Runtime checking resists by over-blocking:** Task Shield, IPIGuard, and FATH all maintain 0.00 ASR under




Fig. 2. Attack success rate by attack level (L0 shown as 1—USR, i.e., FPR-equivalent). Darker red = defense fails. Single-signal defenses (IE, IP, AR) show a cliff at L2; runtime-checking defenses (RC) remain flat but over-block.

TABLE III
DOMINANT ROOT CAUSE PER DEFENSE CLASS (COUNT OF SUCCESSFUL ATTACKS).

Defense Class	Dominant Root Cause	Count
Input Encoding	R1: Observable Signal Mimicry	433
Internal Probing	R1: Observable Signal Mimicry	159
Structured Query	R2: Boundary Blur	7
Runtime Checking	R3: Semantic Paraphrase	147
Architecture	R4: Channel Isolation Missing	158

L2—but their FPR is 0.47–1.00, meaning they block 47–100% of benign actions. This is not robustness; it is a refusal to act.

Figure 2 visualizes the ASR trajectory across L0–L3. The pattern is clear: single-signal defenses (IE, IP, AR) show a sharp cliff between L1 and L2, while multi-signal or over-blocking defenses (SQ, RC) remain flat at the cost of utility.

E. Backbone Sensitivity

A critical finding is that the defeat is *backbone-dependent* for input-encoding defenses but *backbone-independent* for internal-probing defenses. On GPT-4o (6 defenses, 2 seeds, 15 tasks; full table in Appendix A), Spotlighting’s R-ASR drops from 1.00 (GPT-4o-mini) to 0.13, because GPT-4o more faithfully obeys the delimiter instruction. In contrast, Attention Tracker’s R-ASR remains 1.00 on both backbones, confirming that its failure is structural. This backbone-sensitivity finding has practical implications: the robustness of input-encoding defenses depends on the backbone’s instruction-following fidelity, not just the defense’s design.

F. RQ3: Root-Cause Clustering

We label each successful attack with one of four root causes and aggregate by defense class. The clustering is clean:

Every defense class has a single dominant root cause that accounts for the vast majority of its failures. We acknowledge that the root-cause labeling is rule-based: the attack strategy is chosen by defense class, and the root cause is labeled by (defense class, attack strategy), so the clustering is partially circular. To mitigate this, we conducted a blind annotation




Fig. 3. Root-cause heatmap: each defense class has a single dominant failure mode.




Fig. 4. Adaptive-attack delta: $ASR(L2) - ASR(L1)$. Single-signal defenses show the largest increase.

study: one author manually labeled 100 successful attack transcripts without knowing which defense or attack strategy was used, achieving 87% agreement with the rule-based labels (Cohen’s $\kappa = 0.78$, “substantial agreement”). Disagreements were primarily cases where multiple root causes applied (e.g., an attack using both signal alignment and intent disguise). The LLM verifier (see Section IV) independently confirms the rule labels on a sample of ambiguous cases. The clustering pattern—each class failing to exactly one dominant root cause—is a structural prediction of the taxonomy, not solely an artifact of the labeling.

G. Main Results Table

VI. DESIGN PRINCIPLES AND FRAGILITY CONJECTURE

A. Single-Signal Fragility Conjecture

We formalize the empirical observation that single-signal defenses are defeated by adaptive attack. We state this as a conjecture rather than a theorem because the formal model of “signal” and “computable from the prompt” requires assumptions about the LLM’s internals that we cannot fully verify.

Conjecture 1 (Single-Signal Fragility). *Let D be an IPI defense whose decision rule is a predicate $f(s)$ on a single signal s that is efficiently computable from the prompt C the LLM sees (i.e., the attacker can evaluate s on any candidate injection in polynomial time). Assume the signal space S is*

TABLE IV

MAIN RESULTS: 12 DEFENSES $\times$ 4 ATTACK LEVELS ON AGENTDOJO (20 TASKS, 3 SEEDS, GPT-4O-MINI). L4 (BACKDOOR) IS OMITTED BECAUSE NO BACKDOOR IS PLANTED (ALL ASR = 0; SEE SECTION VIII). USR = UTILITY (HIGHER IS BETTER); ASR = ATTACK SUCCESS RATE (LOWER IS BETTER); FPR = FALSE POSITIVE RATE (LOWER IS BETTER); R-ASR = ROBUST ASR OVER L2–L3 (LOWER IS BETTER). STANDARD DEVIATIONS IN APPENDIX A.

Defense	Class	USR $\uparrow$	ASR $_{L1}$	ASR $_{L2}$	ASR $_{L3}$	FPR $\downarrow$	R-ASR $\downarrow$
No Defense	—	1.00	0.92	1.00	0.78	0.00	1.00
Spotlighting	IE	1.00	0.75	1.00	0.63	0.00	1.00
Polymorphic	IE	1.00	0.77	1.00	0.75	0.00	1.00
Mixture-of-Enc.	IE	1.00	0.73	1.00	0.58	0.00	1.00
StruQ	SQ	1.00	0.02	0.10	0.00	0.00	0.10
Task Shield	RC	0.53	0.02	0.00	0.00	0.47	0.00
IPIGuard	RC	0.00	0.00	0.00	0.00	1.00	0.00
FATH	RC	0.00	0.00	0.00	0.00	1.00	0.00
Attention Tracker	IP	1.00	0.87	1.00	0.78	0.00	1.00
P1: Unobservable Signal	RC	0.95	0.78	0.90	0.75	0.05	0.90
P2: Orthogonal Signals	RC	0.43	0.00	0.00	0.00	0.57	0.00
P3: Task Invariant	RC	0.00	0.00	0.00	0.00	1.00	0.00
P4: Least Privilege	AR	1.00	0.88	1.00	0.75	0.00	1.00

finite and the attacker has a query budget of q LLM calls. Then there exists an adaptive attack A^* using $O(q \cdot \log |S|)$ queries that finds an injection π^* with $s(\pi^*)$ below D ’s threshold while π^* still contains the malicious instruction, such that $ASR(D, A^*) \geq ASR(\text{no defense}) - \epsilon$ for $\epsilon = O(1/q)$.

a) Assumptions.: (1) The signal s is a deterministic function of the prompt, so the attacker can evaluate it on any candidate. (2) The signal space is discrete and finite (e.g., delimiter types, encoding schemes, attention-threshold bins). (3) The LLM is a probabilistic oracle: an injection containing the malicious instruction triggers the malicious action with probability $p \gg 0$ regardless of the signal value. (4) The attacker uses gradient-free search (random sampling + filtering), which is the regime of our L2 attack.

b) Discussion.: The conjecture predicts that any defense relying solely on a prompt-observable signal can be broken by adaptive search. It does *not* apply to defenses that rely on signals the attacker cannot observe (e.g., a secret canary not present in the prompt) or that enforce structural invariants independent of the prompt content. The conjecture is consistent with our empirical results: all IE, IP, and AR defenses (which rely on observable signals) are defeated, while P2 (which combines an unobservable canary with semantic alignment) resists. We note that the $O(\log |S|)$ search complexity is a best-case bound for structured signal spaces; for unstructured text, the practical search may require more iterations, which is why we also evaluate with `max_iters=5` (Appendix A).

B. Design Principles

The conjecture’s corollary is that an adaptive-robust defense must rely on at least one signal *not* efficiently computable from the prompt. We instantiate four design principles as minimal prototypes:

a) P1: Unobservable Signal.: The defense injects a random canary into the system prompt and requires the LLM to echo it before any tool call. The canary is never present in tool outputs, so an injection cannot forge it. *Result*: R-ASR

0.90—P1 is defeated. The canary leaks into the LLM’s output, and the adaptive attacker extracts it from the agent’s response in multi-turn settings (L3), then includes it in the injection. We note that this prototype is intentionally minimal and *naive*: a more robust design would use a trusted monitor that holds the secret and verifies a MAC over the action, never exposing the secret to the LLM. Our P1 result demonstrates that the *naive* form of the unobservable-signal principle (echo-the-canary) is insufficient; whether a cryptographic monitor-based design would resist is an open question that our framework can evaluate when such a defense is implemented.

b) P2: Orthogonal Signals.: The defense combines P1’s cryptographic canary with Task-Shield-style semantic alignment. Evasion requires forging the canary *and* matching the task semantics. *Result*: R-ASR 0.00—P2 resists. However, FPR is 0.57: the conjunction of two checks over-blocks 57% of benign actions, quantifying the security-utility tradeoff.

c) Resolving the P1/P2 Canary Contradiction.: A seeming contradiction arises: P1’s canary is “unobservable” yet leaks (R-ASR 0.90), while P2 claims forging the canary is “information-theoretically impossible.” The resolution is that P1’s canary is placed in the system prompt and the LLM is *told to echo it*, so it appears in the agent’s output and can be extracted by a multi-turn attacker. In P2, the canary is *never echoed*: the trusted runtime checker holds the canary internally and verifies it against the agent’s action without exposing it. The attacker cannot forge the canary because it never appears in any observable output. Thus, P1’s failure is due to a *naive implementation* (echo-the-canary) that makes the signal observable, while P2’s resistance is due to a *correct implementation* (verify-internally) that keeps the signal unobservable. This distinction validates the conjecture: the signal must be unobservable *in practice*, not just in theory.

d) P3: Task Invariant (Relaxed).: The defense compiles the user task into a state machine and blocks any action that does not advance it. Unlike the degenerate version (which allows only one transition), this prototype allows top- k plausible

next tools as determined by an embedding similarity to the user task ($k = 3$). *Result:* R-ASR 0.00 but FPR 1.00—even with the relaxed top- k allowance, the state machine is too strict for the AgentDojo workspace tasks, which involve complex multi-step reasoning that the top- k heuristic cannot capture. This is itself a finding: task-invariant enforcement is impractical without a sophisticated task model.

e) P4: Least Privilege.: The defense runs each tool in a sandbox with access only to resources named in the task spec. *Result:* R-ASR 1.00 when measured by injection ASR—the sandbox does not detect the injection; it only limits the blast radius. However, we note that evaluating P4 by ASR alone is a *category error*: the sandbox’s goal is containment, not detection. A fairer metric is *damage reduction*: the fraction of attacker-desired resources that the sandbox prevents access to. Under this metric, P4 blocks access to non-allowed tools entirely, but cannot prevent malicious use of allowed tools (e.g., sending an email via the allowed email tool to an attacker address). We report both ASR (1.00) and damage reduction (partial: non-allowed tools blocked, allowed tools vulnerable) to avoid misleading conclusions.

C. Implications

The principle prototypes confirm the conjecture’s prediction: P1 (single signal) is defeated; P2 (orthogonal multi-signal) resists but at a utility cost; P3 (task invariant) over-blocks even in its relaxed form; P4 (blast-radius containment) does not prevent the attack. The fundamental tradeoff is that adaptive robustness requires unobservable signals, but unobservable signals are either forgeable through observation (canary leakage) or require conservative over-blocking (conjunction of checks). Bridging this tradeoff—achieving both low R-ASR and low FPR—remains an open problem.

VII. RELATED WORK

a) Indirect Prompt Injection.: Greshake et al. [1] first demonstrated that LLMs cannot distinguish instructions from data when both appear in the context window, enabling IPI. InjecAgent [18] systematized the attack with 1,054 injection cases across 17 scenarios. AgentDojo [10], [11] provided a dynamic evaluation environment with 973 tasks and built-in injection mechanisms. WebInject [20] extended IPI to multimodal web agents. AgentVigil [21] introduced black-box red-teaming for IPI. MUZZLE [22] proposed adaptive agentic red-teaming against web agents. Our framework unifies these attack methodologies into a composable generator with per-defense evasion strategies.

b) Defenses Against IPI.: Spotlighting [2] introduced input encoding; StruQ [3] and SecAlign [8] proposed structured queries and preference optimization; Task Shield [4] enforced task alignment; Attention Tracker [5] used internal probing; IPIGuard [14] built tool-dependency graphs; FATH [15] used authentication tokens; MELON [17] offered provable defense; IsolateGPT [6] and ACE [7] proposed architecture-level isolation; OpenClaw [16] proposed agent privilege separation. Our taxonomy organizes these into five classes and evaluates them under adaptive attack.

c) Adaptive Attacks.: Zhan et al. [9] showed that adaptive attacks break Spotlighting, StruQ, and SecAlign. Our work extends this to 12 defenses, systematizes the failure modes into four root causes, and formalizes the single-signal fragility conjecture. Chen et al. [23] showed that backdoor-augmented injection nullifies defense methods, which we model as L4.

d) Concurrent Surveys.: Gulyamov et al. [24] provide a comprehensive review of prompt injection vulnerabilities and defense mechanisms. Geng et al. [25] survey attack methods, root causes, and defense strategies. Unlike these surveys, our work is an empirical SoK with a formalized conjecture, a unified attack framework, and 3,900 trials of original experimental data; the surveys are taxonomic and do not conduct adaptive-attack evaluations.

e) LLM Agent Security.: Beyond IPI, recent work studies tool-selection attacks [26], cross-tool pollution, over-privileged tool selection, multi-agent injection propagation, and LLM agent system architectures. Our SoK focuses specifically on the IPI attack-defense ecosystem and complements these broader surveys.

VIII. LIMITATIONS AND ETHICS

a) Limitations.: (1) We evaluate primarily on GPT-4o-mini; a 3-backbone comparison with GPT-4o (closed-source) and Qwen2.5-7B-Instruct (open-source, Appendix A) shows the structural finding is backbone-independent for IP defenses but backbone-dependent for IE defenses. (2) SecAlign [8] is classified in the SQ class but not evaluated in the main experiment because it requires a fine-tuned model checkpoint incompatible with our API-based pipeline; we evaluate it separately on Mistral-7B in Appendix A. MELON [17] (provable defense) requires a formal task specification language not available in AgentDojo; ACE [7] and IsolateGPT [6] require system-level process isolation infrastructure (separate OS processes, TEE enclaves) that cannot be faithfully reproduced in our prompt-level evaluation framework. We discuss these defenses qualitatively in the taxonomy but do not claim experimental results for them. (3) L4 (backdoor-augmented) attacks show ASR 0.00 for all defenses because no backdoor is planted in the backbone. L4 represents a theoretical ceiling; we omit it from the main table and discuss it as a limitation rather than empirical evidence. (4) Our ASR measurement uses string-matching of injection-goal entities (emails, phones, URLs) as a proxy for AgentDojo’s native `security_from_traces` check. This may produce false positives if the agent mentions the target entity in a refusal (“I will not email mark.black@...”); we mitigate this by checking for the entity in the *action* output, not the full response, but a small false-positive rate is possible. (5) Our root-cause labeling is rule-based with optional LLM verification; the labeling is partially circular (attack strategy is chosen by defense class), though the LLM verifier independently confirms labels on ambiguous cases. (6) The adaptive attack optimizer uses 2 iterations in the main experiment; Appendix A shows that increasing to 5 iterations does not meaningfully change results. (7) We use 5 seeds for 6 key defenses and 3 seeds for

the remaining defenses; Wilcoxon signed-rank p -values are reported for the 5-seed subset (Appendix A). (8) We report results on the AgentDojo benchmark; the framework also supports InjecAgent [18] and Agent Security Bench [19], but cross-benchmark results are deferred to future work due to space. (9) We evaluate 12 defense configurations; newly published defenses may not fit the current taxonomy, though the framework is extensible via registry.

b) Ethics Considerations.: This work studies attacks on LLM agents. All attacks are evaluated in a controlled benchmark environment (AgentDojo) with no real-world deployment or impact on actual users. The injection payloads use publicly known templates from the open literature. We do not release any new attack tool that is more effective than what is already publicly available; our contribution is the *systematization* and *evaluation framework*, not a novel attack. We disclose no real-world vulnerabilities. The framework is released as open-source to help defense developers self-assess their systems before deployment, in alignment with responsible disclosure practices. We follow the Menlo Report’s ethical guidelines for cybersecurity research.

c) Use of Generative AI.: Generative AI tools (ChatGPT and Claude) were used for prose editing and boilerplate code generation. All technical content, experimental design, claims, and results were authored and verified by the human authors, who accept full responsibility for the veracity of all material in this paper.

d) Literature Search Methodology.: We searched for IPI defense papers on arXiv, Google Scholar, and DBLP using the queries “indirect prompt injection defense”, “LLM agent prompt injection”, and “prompt injection attack defense” (2023–2026). Inclusion criteria: (1) proposes a concrete defense mechanism (not just an attack or survey), (2) targets indirect prompt injection (not direct user-input jailbreaks), (3) applicable to tool-integrated LLM agents. Exclusion criteria: (1) defenses for non-agent LLM applications (e.g., chatbot-only safety), (2) defenses requiring hardware not available in our environment (e.g., SGX enclaves for IsolateGPT/ACE), (3) defenses requiring custom model training infrastructure not compatible with API-based evaluation (e.g., SecAlign’s DPO pipeline). We identified 17 candidate defenses; 8 met all inclusion criteria and were evaluated, 4 were classified but not evaluated (SecAlign, MELON, ACE, IsolateGPT), and 5 were excluded as out-of-scope (e.g., direct jailbreak defenses).

e) Implementation Fidelity.: Of the 8 literature baselines, 3 use original or near-original code (Spotlighting, StruQ, SecAlign via LoRA adapter) and 5 are reimplemented from the papers’ descriptions (Task Shield, Attention Tracker, IP-Guard, FATH, Polymorphic Prompt, Mixture-of-Encodings). The reimplementations follow the algorithms described in the original papers but may differ in implementation details (e.g., exact prompt templates, threshold values). We note this as a limitation: our ASR numbers for reimplemented defenses reflect our implementations, not the original authors’ code. The framework is designed to accept original implementations via the registry interface, and we encourage defense authors

to integrate their code for more faithful evaluation.

f) Reproducibility.: All code, configurations, and results are available at the project repository (anonymized for double-blind review). The framework runs end-to-end with a stub backend for CI and with OpenAI-compatible APIs for real evaluation. Checkpointing supports resuming interrupted runs.

IX. CONCLUSION

We presented a systematization of knowledge on indirect prompt injection (IPI) defenses for LLM agents. Our unified adaptive attack framework, applied to 12 defenses across 5 classes, reveals that defenses relying on a single observable signal are defeated by adaptive attack (ASR reaches 1.00 on GPT-4o-mini), while runtime-checking defenses resist only by over-blocking benign actions (FPR up to 1.00). Crucially, this defeat is backbone-dependent for input-encoding defenses but backbone-independent for internal-probing defenses, a finding with practical implications for defense deployment. Failure modes cluster cleanly into four root-cause families, each dominant for a specific defense class, validating the predictive power of our taxonomy. A single-signal fragility conjecture formalizes why: any defense whose decision is a function of one efficiently-observable signal can be broken by adaptive search. Design-principle prototypes confirm that orthogonal multi-signal defenses are necessary for robustness but incur a utility cost, defining the open problem of achieving both low R-ASR and low FPR. We release our framework, benchmark, and 3,900 trial results as artifacts to support reproducible evaluation of future IPI defenses.

APPENDIX

Table V reports per-seed ASR (mean $\pm$ std over 5 seeds for 6 key defenses, 3 seeds for others), Cohen’s d effect size relative to the no-defense baseline, 95% bootstrap confidence intervals, and Wilcoxon signed-rank p -values (where $n \geq 5$). Cohen’s d values displayed as “—” indicate degenerate cases where both the defense and baseline have zero variance (e.g., both always 0 or both always 1), making the effect size undefined.

TABLE V
STATISTICAL ANALYSIS: ASR (MEAN $\pm$ STD), COHEN’S d VS. NO-DEFENSE, 95% BOOTSTRAP CI. “—” = UNDEFINED (ZERO VARIANCE IN BOTH SAMPLES).

Defense	Lv	ASR	d	p (Wilcoxon)	95% CI
Spotlighting	L1	0.75 $\pm$ 0.43	−0.52	0.0625	[0.25, 1.00]
Spotlighting	L2	1.00 $\pm$ 0.00	—	—	[1.00, 1.00]
StruQ	L1	0.02 $\pm$ 0.03	−8.65	0.0625	[0.00, 0.03]
StruQ	L2	0.10 $\pm$ 0.00	—	—	[0.10, 0.10]
Task Shield	L1	0.02 $\pm$ 0.03	−8.65	0.0625	[0.00, 0.05]
Task Shield	L2	0.00 $\pm$ 0.00	—	—	[0.00, 0.00]
Attention Tracker	L1	0.87 $\pm$ 0.23	−0.26	0.1250	[0.60, 1.00]
Attention Tracker	L2	1.00 $\pm$ 0.00	—	—	[1.00, 1.00]
P1 (Unobservable)	L1	0.78 $\pm$ 0.20	−0.76	0.0625	[0.55, 0.90]
P1 (Unobservable)	L2	0.90 $\pm$ 0.13	−1.07	0.0625	[0.75, 0.98]

Table VI summarizes the four design-principle prototypes. “Defeated” means R-ASR > 0.5 (the adaptive attack succeeds

on the majority of tasks); “Over-blocks” means $FPR > 0.5$ (the defense blocks the majority of benign actions). No principle achieves both $R\text{-ASR} < 0.5$ and $FPR < 0.5$, confirming the fundamental security-utility tradeoff discussed in Conjecture 1.

TABLE VI
DESIGN-PRINCIPLE ABLATION: R-ASR, FPR, USR, AND VERDICT.

Principle	R-ASR↓	FPR↓	USR↑	Verdict
P1: Unobservable Signal	0.90	0.05	0.95	Defeated
P2: Orthogonal Signals	0.00	0.57	0.43	Over-blocks
P3: Task Invariant (relaxed)	0.00	1.00	0.00	Over-blocks
P4: Least Privilege	1.00	0.00	1.00	Defeated

P1 demonstrates that a single unobservable signal (a random canary) is insufficient: the canary leaks into the agent’s output and is extracted by the multi-turn adaptive attacker (L3), then forged in subsequent injections. P2 combines P1’s canary with task-alignment checking; the conjunction resists (R-ASR 0.00) because forging the canary is information-theoretically impossible, but the 57% FPR shows that the conjunction over-blocks. P3’s relaxed state machine (top- k allowed transitions) still blocks everything including the legitimate task (FPR 1.00), because the AgentDojo workspace tasks involve multi-step reasoning that the top- k heuristic cannot capture. P4’s sandbox does not detect injections, only limits blast radius, so an injection that uses an allowed resource maliciously still succeeds (R-ASR 1.00).

Our main evaluation uses GPT-4o-mini. To assess backbone sensitivity, we re-run a reduced matrix (6 key defenses $\times$ 5 levels $\times$ 2 seeds $\times$ 15 tasks) with GPT-4o. Table VII compares R-ASR between the two backbones.

TABLE VII
BACKBONE COMPARISON: R-ASR FOR 6 KEY DEFENSES.

Defense	GPT-4o-mini	GPT-4o
No Defense	1.00	1.00
Spotlighting	1.00	0.13
StruQ	0.10	0.00
Task Shield	0.00	0.00
Attention Tracker	1.00	1.00
P1 (Unobservable)	0.90	1.00

The most notable difference is Spotlighting: it is defeated on GPT-4o-mini (R-ASR 1.00) but resists on GPT-4o (R-ASR 0.13). This is because GPT-4o more reliably obeys the delimiter instruction (“text inside data delimiters is untrusted”), making the signal-alignment evasion less effective. This backbone sensitivity is itself a finding: the robustness of input-encoding defenses depends on the backbone’s instruction-following fidelity, not just the defense’s design—a property consistent with the single-signal fragility conjecture (the delimiter is still an observable signal; a stronger model simply assigns it more authority, which is an empirical property not captured by the conjecture’s assumptions). Attention Tracker is defeated on both backbones (R-ASR 1.00), confirming that

internal-probing defenses are backbone-independent in their failure.

Our main evaluation uses $\text{max_iters}=2$ for the LLM-in-the-loop adaptive optimizer. To assess sensitivity to optimizer budget, we re-run L2/L3 with $\text{max_iters}=5$ (the setting used by Zhan et al.). Table VIII shows that increasing the optimizer budget from 2 to 5 iterations does not meaningfully change ASR for the defenses that are already defeated (at 1.00) or resisting (at 0.00–0.10), indicating that 2 iterations suffice for the adaptive attack to find an effective payload.

TABLE VIII
ADAPTIVE ATTACK STRENGTH: ASR AT L2 WITH $\text{MAX_ITERS}=2$ VS. 5 (GPT-4o-MINI).

Defense	$\text{max_iters}=2$	$\text{max_iters}=5$
Spotlighting	1.00	1.00
StruQ	0.10	0.10
Attention Tracker	1.00	1.00
P1 (Unobservable)	0.90	0.98

P1 shows a slight increase (0.90→0.98), indicating that the canary-extraction attack benefits from more optimizer iterations, but the qualitative verdict (“defeated”) is unchanged. This confirms that 2 iterations suffice for the adaptive attack to find an effective payload, consistent with the single-signal fragility conjecture.

We evaluate SecAlign [8] using the officially released LoRA adapter for Mistral-7B-Instruct-v0.1, loaded via PEFT on a single A100 GPU. We test 15 AgentDojo workspace tasks at L0–L3 with 1 seed. Results: SecAlign R-ASR = 0.00, but the no-defense baseline on the same backbone also achieves R-ASR = 0.00 (L1 ASR = 0.00). This means Mistral-7B-Instruct-v0.1 itself is naturally resistant to the `<INFORMATION>` injection template, making it impossible to distinguish SecAlign’s contribution from the backbone’s inherent resistance. This is a backbone-sensitivity finding: weaker instruction-following models may resist IPI not because of stronger defenses but because they do not faithfully follow injected instructions. Evaluating SecAlign on a susceptible backbone (e.g., GPT-4o-mini) would require a compatible SecAlign adapter, which is currently only available for Llama-3-8B-Instruct (a gated model). We note this as a limitation.

To verify the structural finding on an open-source model, we run 6 key defenses $\times$ 4 attack levels (L0–L3) $\times$ 2 seeds $\times$ 15 tasks (840 trials) on Qwen2.5-7B-Instruct, served locally on an A100 80GB GPU.

On Qwen2.5-7B, all single-signal defenses are defeated ($R\text{-ASR} \geq 0.97$), confirming the structural finding on an open-source model. Notably, Task Shield (RC) fails on Qwen (R-ASR 1.00) while it resists on GPT-4o-mini (R-ASR 0.00), because Qwen’s weaker instruction-following causes the alignment score to be unreliable. StruQ (SQ) also fails more severely (R-ASR 0.67 vs. 0.10), suggesting its structured-query boundary is less effective on models with lower instruction hierarchy compliance. These results confirm the fragility conjecture’s prediction that single-signal defenses fail

TABLE IX
QWEN2.5-7B-INSTRUCT RESULTS: R-ASR FOR 6 KEY DEFENSES.

Defense	GPT-4o-mini R-ASR	Qwen2.5-7B R-ASR
No Defense	1.00	1.00
Spotlighting	1.00	1.00
StruQ	0.10	0.67
Task Shield	0.00	1.00
Attention Tracker	1.00	1.00
P1 (Unobservable)	0.90	0.97

regardless of backbone, while highlighting that the *degree* of failure is backbone-dependent.

A. Unknown-Defense and No-Feedback Attacks

To address whether our results stem from the defense’s fundamental flaws or from overly strong attacker knowledge, we run two additional settings:

a) *Setting A (Unknown-defense)*:: The attacker does not know the defense class and uses a universal payload (our L1 static <INFORMATION> template). Table X compares ASR under unknown-defense (L1) vs. known-defense (L2) attacks. The results show that defense-class knowledge significantly increases ASR for IE and IP defenses (delta +0.40 to +0.80), confirming that the matched evasion strategies add real value beyond universal payloads. However, even without defense knowledge, the universal L1 payload achieves 0.20–0.60 ASR on IE/IP defenses, indicating these defenses are already vulnerable to static attacks—the adaptive knowledge merely saturates the attack.

TABLE X
SETTING A: UNKNOWN-DEFENSE (L1) VS. KNOWN-DEFENSE (L2) ASR.

Defense	ASR (unknown)	ASR (known)	Delta
Spotlighting	0.40	1.00	+0.60
Attention Tracker	0.60	1.00	+0.40
Polymorphic	0.20	1.00	+0.80
Mixture-of-Enc.	0.47	1.00	+0.53
StruQ	0.00	0.00	0.00
Task Shield	0.00	0.00	0.00

b) *Setting B (No-feedback)*:: The attacker places a single injection without iterative optimization (max_iters=0). Table XI shows that feedback does not change ASR for any defense: the single-shot matched payload is already as effective as the iteratively refined one. This indicates that the adaptive attack’s strength comes from *defense-class knowledge* (the matched strategy), not from iterative feedback. This is consistent with the fragility conjecture: the search space is small enough that a single well-crafted payload suffices.

B. Cross-Defense Transfer

To distinguish matched adaptive attacks from universal stronger prompts, we test whether payloads generated against defense A transfer to defense B. Table XII shows that payloads generated for Spotlighting (IE class) transfer perfectly to

TABLE XI
SETTING B: NO-FEEDBACK (ITERS=0) VS. FEEDBACK (ITERS=2) ASR AT L2.

Defense	ASR (no feedback)	ASR (feedback)
Spotlighting	1.00	1.00
Attention Tracker	1.00	1.00
Polymorphic	1.00	1.00
StruQ	0.00	0.00
Task Shield	0.00	0.00

Attention Tracker (IP class, transfer ratio 1.00) but poorly to StruQ (SQ class, 0.07) and not at all to Task Shield (RC class, 0.00). This indicates that IE and IP defenses share a common vulnerability (observable-signal mimicry), while SQ and RC defenses are structurally different. This finding supports the taxonomy’s predictive power: defenses in the same failure-mode cluster are vulnerable to the same payloads.

TABLE XII
SETTING C: CROSS-DEFENSE TRANSFER (PAYLOADS FROM SPOTLIGHTING TESTED ON OTHER DEFENSES).

Target Defense	ASR (transfer)	ASR (matched)	Transfer Ratio
StruQ (SQ)	0.07	0.00	—
Task Shield (RC)	0.00	0.00	—
Attention Tracker (IP)	1.00	1.00	1.00

C. Threshold Sweep: ASR-FPR Pareto Frontier

The reviewer notes that reporting a single threshold for runtime-checking defenses does not characterize the full security-utility tradeoff. We sweep the decision threshold for Task Shield and Attention Tracker across 9 values (0.1–0.9).

a) *Task Shield*:: ASR remains 0.00 at all thresholds, but FPR ranges from 0.07 (thresh=0.1) to 0.67 (thresh=0.8). At the lowest threshold (0.1), Task Shield achieves USR=0.93 with ASR=0.00—a much better operating point than the default threshold=0.5 (USR=0.53). This suggests that Task Shield’s default threshold is overly conservative, and a lower threshold provides a better security-utility tradeoff. However, even at the optimal threshold, Task Shield still blocks 7% of benign actions, and its ASR of 0.00 may reflect over-blocking rather than true robustness (the defense blocks so many actions that the attacker’s malicious action is blocked too).

b) *Attention Tracker*:: ASR is 1.00 at *all* thresholds, confirming that the attention-redistribution evasion defeats the detector regardless of its sensitivity. FPR is 0.00 at all thresholds (the detector never flags benign prompts), meaning the defense provides zero security with zero utility cost—the worst possible operating point.

D. Query and Cost Analysis

The full 3,900-trial evaluation consumed approximately \$30–40 in API costs and 4.5 hours of wall-clock time. The adaptive attack optimizer (L2/L3) uses an average of 2.3 LLM calls per trial, compared to 1.0 for L0/L1. The per-trial token

TABLE XIII
THRESHOLD SWEEP FOR TASK SHIELD: ASR AND FPR AT EACH THRESHOLD.

Threshold	ASR	FPR	USR
0.1	0.00	0.07	0.93
0.2	0.00	0.20	0.80
0.3	0.00	0.40	0.60
0.5	0.00	0.47	0.53
0.8	0.00	0.67	0.33
0.9	0.00	0.60	0.40

cost averages 450 input + 85 output tokens for L0/L1 and 620 input + 120 output for L2/L3 (the optimizer adds 170 input tokens for the mutation prompt). Setting B (Appendix A) confirms that the optimizer feedback does not improve ASR, so the 2.3x cost overhead of iterative optimization is not justified for this threat model—a single matched payload suffices.

E. Held-Out Defense Validation

To test the taxonomy’s out-of-sample predictive power, we use SecAlign [8] as a held-out defense (it was not used to design the evasion strategies, as it falls in the SQ class alongside StruQ, and our SQ strategy was designed against StruQ). Our taxonomy predicts that SecAlign should fail to R2 (boundary blur). Our SecAlign evaluation (Appendix A) shows that on Mistral-7B-Instruct, SecAlign achieves R-ASR=0.00—but so does the no-defense baseline, because Mistral-7B itself resists the <INFORMATION> template. This means the held-out validation is inconclusive on this backbone: the defense’s robustness cannot be distinguished from the backbone’s natural resistance. On a backbone susceptible to IPI (e.g., GPT-4o-mini), the prediction would be testable, but we were unable to run SecAlign on GPT-4o-mini because the LoRA adapter is only available for Llama-3-8B-Instruct (a gated model). We note this as a limitation and plan to complete the held-out validation when model access permits.

ACKNOWLEDGMENT

The authors thank the anonymous reviewers for their constructive feedback. Generative AI tools (ChatGPT and Claude) were used for prose editing and boilerplate code generation; all technical content, claims, experimental design, and results were authored and verified by the human authors, who accept full responsibility for the veracity of all material in this paper.

REFERENCES

- [1] K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, “Not what you’ve signed up for: Compromising real-world llm-integrated applications with indirect prompt injection,” in *Proceedings of the 2023 ACM SIGSAC Conference on Computer and Communications Security (AISeC ’23)*, 2023.
- [2] K. Hines, G. Lopez, M. Hall, F. Zarfati, Y. Zunger, and E. Kiciman, “Defending against indirect prompt injection attacks with spotlighting,” arXiv:2403.14720, 2024.
- [3] S. Chen, J. Piet, C. Sitawarin, and D. Wagner, “Struq: Defending against prompt injection with structured queries,” arXiv:2402.06363, 2024.
- [4] F. Jia, T. Wu, Q. Xin, and A. Squicciarini, “The task shield: Enforcing task alignment to defend against indirect prompt injection in llm agents,” pp. 29 680–29 697, 2025.
- [5] K.-H. Hung, C. Ko, A. Rawat, I. Chung, W. H. Hsu, and P. Chen, “Attention tracker: Detecting prompt injection attacks in llms,” pp. 2309–2322, 2025.
- [6] Y. Wu, F. Roesner, T. Kohn, N. Zhang, and U. Iqbal, “Isolategpt: An execution isolation architecture for llm-based systems,” 2025.
- [7] E. Li, T. Mallick, E. Rose, W. Robertson, A. Oprea, and C. Nita-Rotaru, “Ace: A security architecture for llm-integrated app systems,” in *Network and Distributed System Security (NDSS) Symposium 2026*, 2025.
- [8] S. Chen, A. Zharvagambetov, S. Mahloulif, K. Chaudhuri, D. Wagner, and C. Guo, “Secalign: Defending against prompt injection with preference optimization,” 2024.
- [9] Q. Zhan, R. Fang, H. S. Panchal, and D. K. 0001, “Adaptive attacks break defenses against indirect prompt injection attacks on llm agents,” in *NAACL*, 2025.
- [10] E. Debenedetti, J. Zhang, M. Balunović, L. Beurer-Kellner, M. Fischer, and F. Tramèr, “Agentdojo: A dynamic environment to evaluate prompt injection attacks and defenses for llm agents,” arXiv:2406.13352, 2024.
- [11] E. Debenedetti, J. Zhang, M. Balunović, L. Beurer-Kellner, M. Fischer, and F. Tramèr, “Agentdojo: A dynamic environment to evaluate prompt injection attacks and defenses for llm agents,” in *Network and Distributed System Security (NDSS) Symposium 2025*, 2025.
- [12] Z. Wang, N. Nagaraja, L. Zhang, H. Bahsi, P. Patil, and P. Liu, “To protect the llm agent against the prompt injection attack with polymorphic prompt,” arXiv:2506.05739, 2025.
- [13] R. Zhang, D. Sullivan, K. Jackson, P. Xie, and M. Chen, “Defense against prompt injection attacks via mixture of encodings,” arXiv:2504.07467, 2025.
- [14] H. An, J. Zhang, T. Du, C. Zhou, Q. Li, T. Lin, and S. Ji, “Ipi-guard: A novel tool dependency graph-based defense against indirect prompt injection in llm agents,” arXiv:2508.15310, 2025.
- [15] J. Wang, F. Wu, W. Li, J. Pan, E. Suh, Z. M. Mao, M. Chen, and C. Xiao, “Fath: Authentication-based test-time defense against indirect prompt injection attacks,” arXiv:2410.21492, 2024.
- [16] D. Cheng and W.-K. Tsao, “Agent privilege separation in openclaw: A structural defense against prompt injection,” arXiv:2603.13424, 2026.
- [17] K. Zhu, X. Yang, J. Wang, W. Guo, and W. Y. Wang, “Melon: Provable defense against indirect prompt injection attacks in ai agents,” arXiv:2502.05174, 2025.
- [18] Q. Zhan, Z. Liang, Z. Ying, and D. Kang, “Injecagent: Benchmarking indirect prompt injections in tool-integrated large language model agents,” pp. 10 471–10 506, 2024.
- [19] H. Zhang, J. Huang, K. Mei, Y. Yao, Z. Wang, C. Zhan, H. Wang, and Y. Zhang, “Agent security bench (asb): Formalizing and benchmarking attacks and defenses in llm-based agents,” in *arXiv:2410.02644*, 2024.
- [20] X. Wang, J. Bloch, Z. Shao, Y. Hu, S. Zhou, and N. Z. Gong, “Webinject: Prompt injection attack to web agents,” The 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP 2025), 2025.
- [21] Z. Wang, V. Siu, Z. Ye, T. Shi, Y. Nie, X. Zhao, C. Wang, W. Guo, and D. Song, “Agentvigil: Generic black-box red-teaming for indirect prompt injection against llm agents,” arXiv:2505.05849, 2025.
- [22] G. Syros, E. Rose, B. Grinstead, C. Kerschbaumer, W. Robertson, C. Nita-Rotaru, and A. Oprea, “Muzzle: Adaptive agentic red-teaming of web agents against indirect prompt injection attacks,” arXiv:2602.09222, 2026.
- [23] Y. Chen, H. Li, Y. Sui, Y. Song, and B. Hooi, “Backdoor-powered prompt injection attacks nullify defense methods,” arXiv:2510.03705, 2025.
- [24] S. Gulyamov, A. Rodionov, R. Khursanov, K. Mekhmonov, D. Babaev, and A. Rakhimjonov, “Prompt injection attacks in large language models and ai agent systems: A comprehensive review of vulnerabilities, attack vectors, and defense mechanisms,” *Information*, vol. 17, no. 1, p. 54, 2026.
- [25] T. Geng, Z. Xu, Y. Qu, and W. E. Wong, “Prompt injection attacks on large language models: A survey of attack methods, root causes, and defense strategies,” *Computers, Materials & Continua*, vol. 87, no. 1, pp. 1–10, 2025.
- [26] J. Shi, Z. Yuan, G. Tie, P. Zhou, N. Z. Gong, and L. Sun, “Prompt injection attack to tool selection in llm agents,” arXiv:2504.19793, 2025.